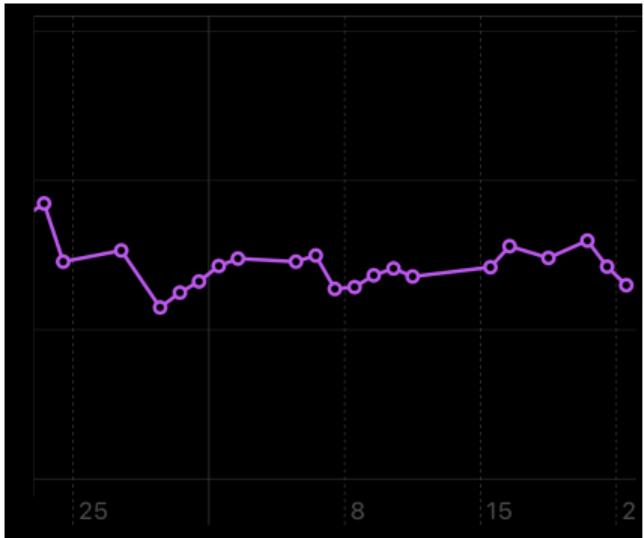


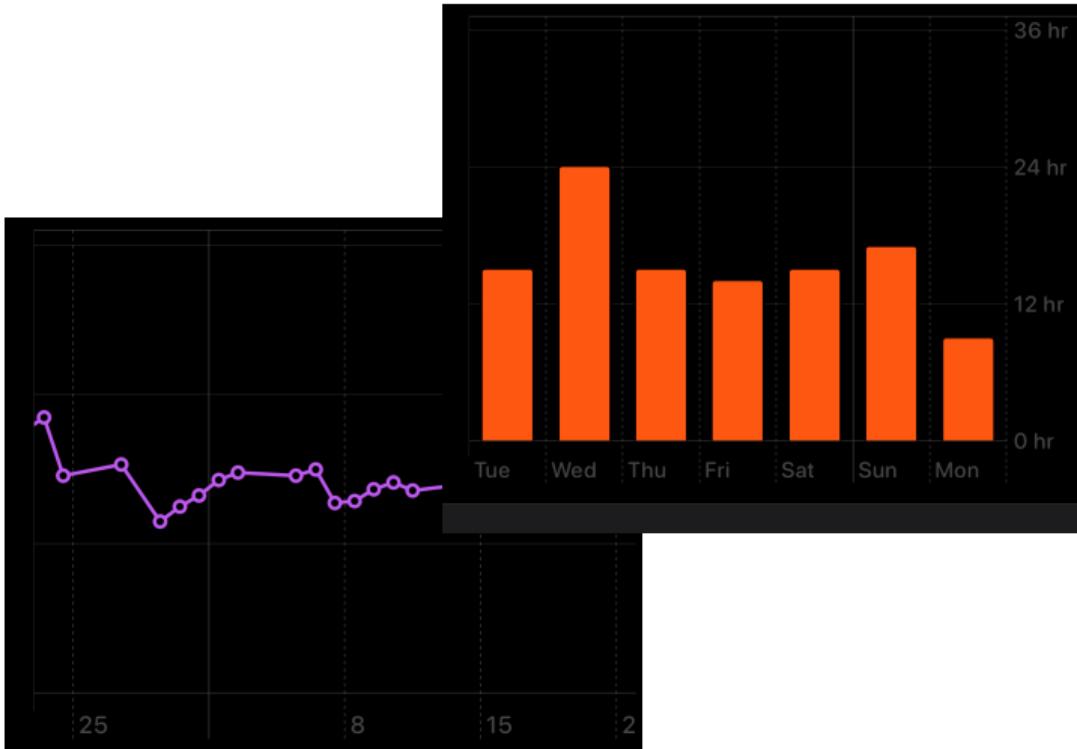
FRUSTRATIONS OF REDUCTIVE SYSTEMS



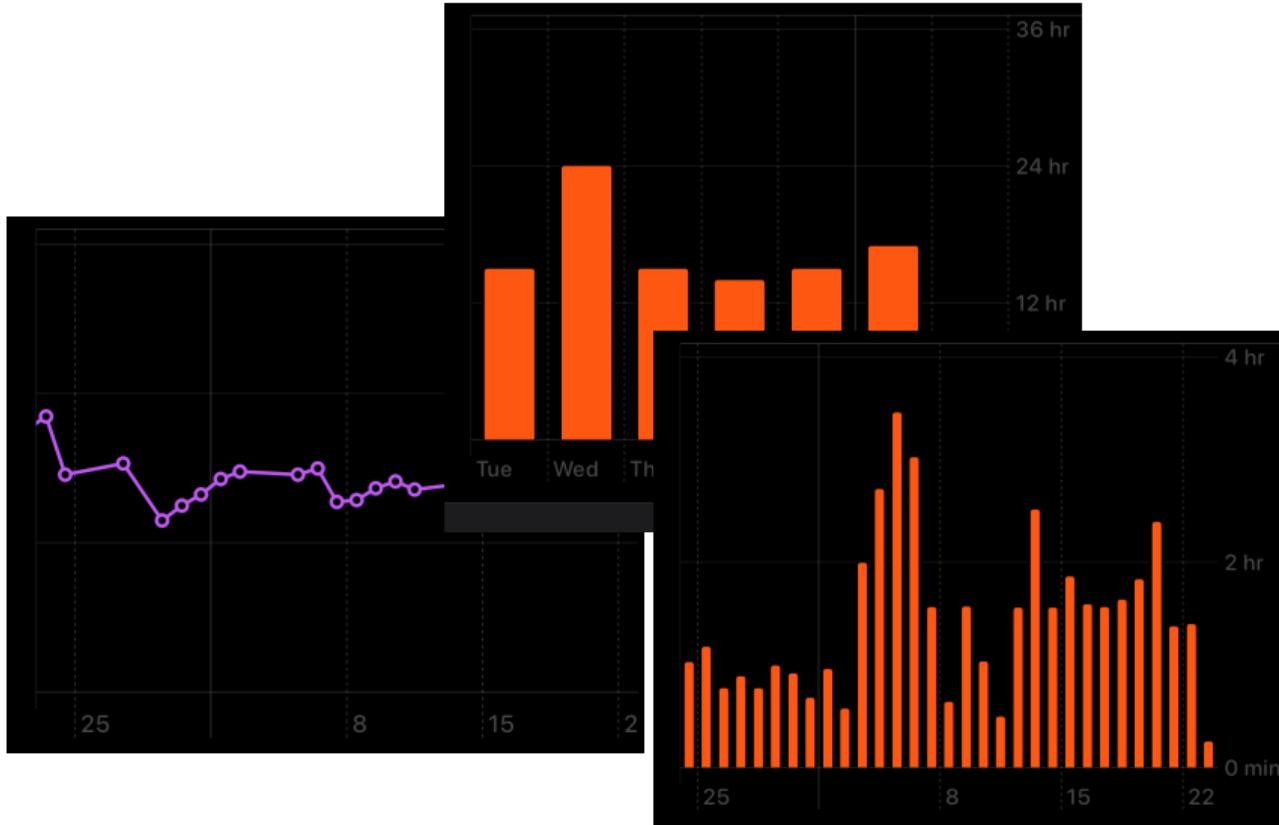
DANGERS OF “THICK DESCRIPTION”



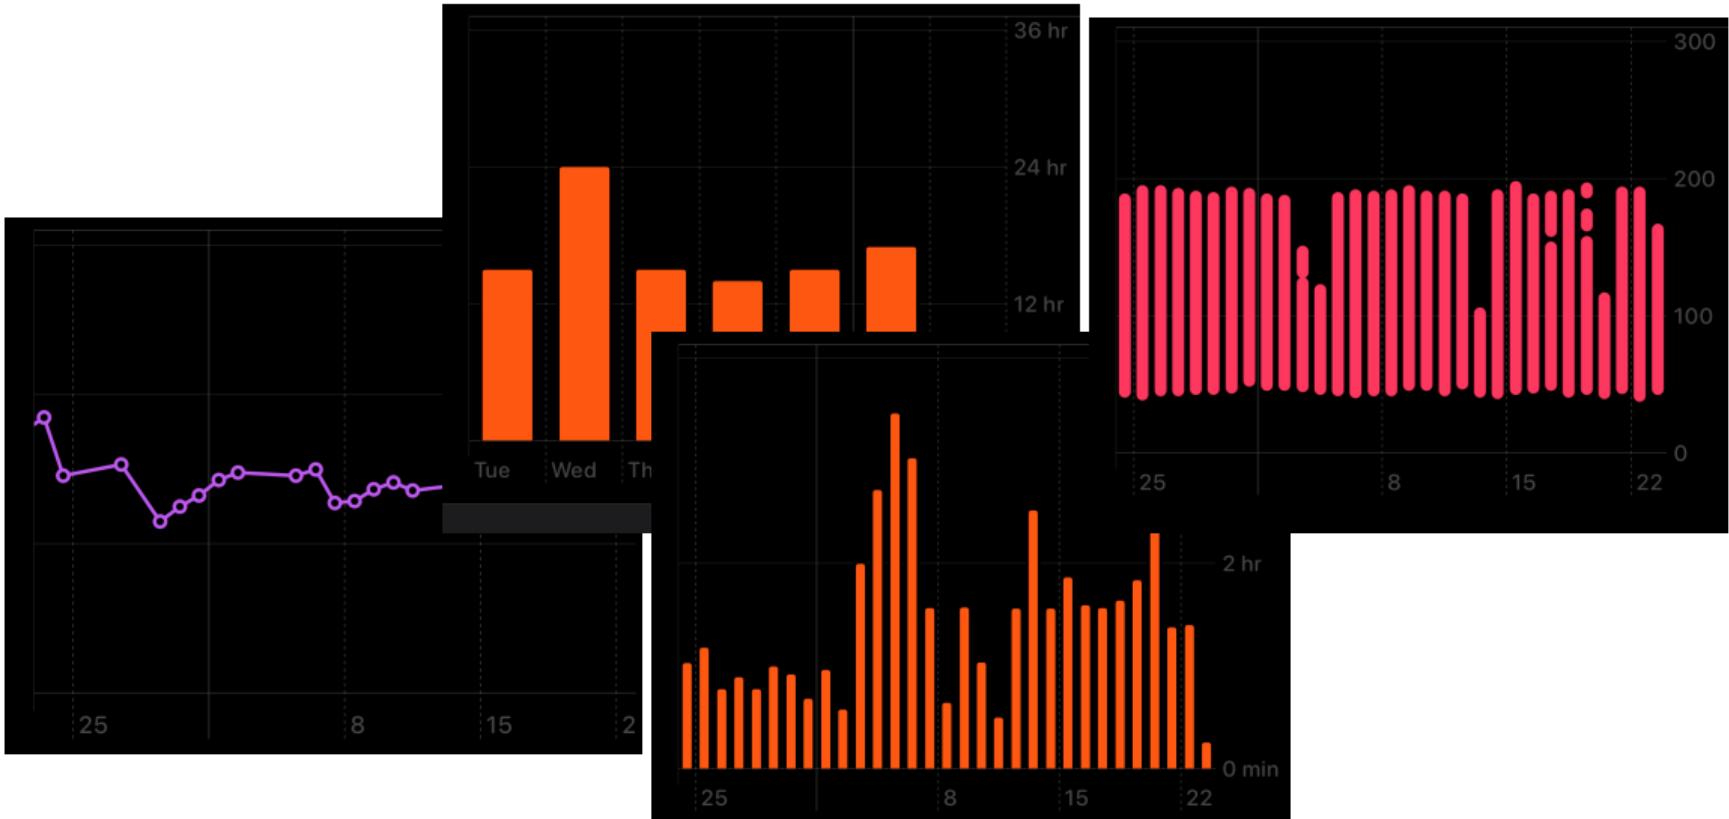
DANGERS OF “THICK DESCRIPTION”



DANGERS OF “THICK DESCRIPTION”



DANGERS OF “THICK DESCRIPTION”



Emotion Regulation for Frustrating Driving Contexts

Helen Harris

Stanford University, Dept. of Communication
450 Serra Mall, Building 120
Stanford, CA. 94305-2050
helenh@stanford.edu

Clifford Nass

Stanford University, Dept. of Communication
450 Serra Mall, Building 120
Stanford, CA. 94305-2050
nass@stanford.edu

ABSTRACT

Driving is a challenging task because of the physical, attentional, and emotional demands. When drivers become frustrated by events their negative emotional state can escalate dangerously. This study examines behavioral and attitudinal effects of cognitively reframing frustrating events. Participants ($N = 36$) were asked to navigate a challenging driving course that included frustrating events such as long lights and being cut-off. Drivers were randomly assigned to three conditions. After encountering a frustrating event, drivers in a *reappraisal-down* condition heard voice prompts that reappraised the event in an effort to deflate negative reactions. Drivers in the second group, *reappraisal-up*, heard voice prompts that brought attention to the negative actions of vehicles and pedestrians. Drivers in a *silent* condition drove without hearing any voice

if their emotional state is known [9], near-term solutions should use available knowledge of the road to anticipate driver frustrations.

Emotion Regulation

The field of psychology provides a significant body of work to aid in addressing negative emotions and promoting healthier states. The process model of emotion regulation [3, 5] posits that emotions may be regulated at one of five points during the time course of emotion: selection of the situation; modification of the situation; deployment of attention; change of cognitions; and modulation of the response. If we aim to improve frustration during everyday circumstances, not all of these stages are feasible points for an intervention.

Considering the selection of the situation, drivers can

Emotion Regulat

Helen Hart

Stanford University, Dept. of
Computer Science
450 Serra Mall, Building 50
Stanford, CA, 94305
helenh@stanford.edu

ABSTRACT

Driving is a challenging task because it involves attentional, cognitive, and emotional demands. When drivers are frustrated by events their negative emotions can escalate dangerously. This study examined the cognitive and emotional effects of cognitively demanding events. Participants ($N = 36$) were assigned to a challenging driving course that included such as long lights and being caught in traffic. Participants were randomly assigned to three conditions: a control condition, a condition where they experienced a frustrating event, or a condition where they heard voice prompts that reappraised the event to deflate negative reactions. Drive times were measured and participants completed a questionnaire about their emotional state.

CHI 2019 Paper

CHI 2019, May 4–9, 2019, Glasgow, Scotland, UK

Using Time and Space Efficiently in Driverless Cars: Findings of a Co-Design Study

Gunnar Stevens
Information Systems
University of Siegen
Siegen, Germany
gunnar.stevens@uni-siegen.de

Paul Bossauer
Department of Management Sciences
Bonn Rhein-Sieg University
Sankt Augustin, Germany
paul.bossauer@h-brs.de

Stephanie Vonholdt
Department of Management Sciences
Bonn-Rhein-Sieg University
Sankt Augustin, Germany
stephanie.vonholdt@h-brs.de

Christina Pakusch
Department of Management Sciences
Bonn Rhein-Sieg University
Sankt Augustin, Germany
christina.pakusch@h-brs.de

ABSTRACT

The alternative use of travel time is one of the widely discussed benefits of driverless cars. We therefore conducted a co-design study with driverless car manufacturers and passengers to explore how driverless cars could be used to make better use of travel time.

Design Study. In *2019 CHI Conference on Human Factors in Computing Systems Proceedings (CHI 2019)*, May 4–9, 2019, Glasgow, Scotland, UK. ACM, New York, NY, USA. Paper 405, 11 pages. <https://doi.org/10.1145/3290607.3300635> 1234567890





CHI 2019 Paper



CHI 2019, May 4–9, 2019, Glasgow, Scotland, UK

Unremarkable AI: Fitting Intelligent Decision Support into Critical, Clinical Decision-Making Processes

Qian Yang

HCI Institute

Carnegie Mellon University

yangqian@cmu.edu

Aaron Steinfeld

Robotics Institute

Carnegie Mellon University

steinfeld@cmu.edu

John Zimmerman

HCI Institute

Carnegie Mellon University

johnz@cs.cmu.edu

ABSTRACT

Clinical decision support tools (DST) promise improved healthcare outcomes by offering data-driven insights. While effective in lab settings, almost all DSTs have failed in practice. Empirical research diagnosed poor contextual fit as the cause. This paper describes the design and field evaluation of a radically new form of DST. It automatically generates slides for clinicians' decision meetings with subtly embedded machine prognostics. This design took inspiration from the notion of *Unremarkable Computing*, that by augmenting the users' routines technology/AI can have significant importance for the users yet remain unobtrusive. Our field evaluation suggests clinicians are more likely to encounter and embrace such a DST. Drawing on their responses, we discuss the importance and intricacies of finding the right level of unremarkability in DST design, and share lessons learned in prototyping critical AI systems as a situated experience.

CCS CONCEPTS

• Human-centered computing → User centered design;

KEYWORDS

1 INTRODUCTION

The idea of leveraging machine intelligence in healthcare in the form of decision support tools (DSTs) has fascinated healthcare and AI researchers for decades. These tools often promise insights on patient diagnosis, treatment options, and likely prognosis. With the adoption of electronic medical records and the explosive technical advances in machine learning (ML) in recent years, now seems a perfect time for DSTs to impact healthcare practice.

Interestingly, almost all these tools have failed when migrating from research labs to clinical practice in the past 30 years [5, 8, 9]. In a review of deployed DSTs, healthcare researchers ranked the lack of HCI considerations as the most likely reason for failure [12, 23]. This includes a lack of consideration for clinicians' workflow and the collaborative nature of clinical work. The interaction design of most clinical decision support tools instead assumes that individual clinicians will recognize when they need help, walk up and use a system that is separate from the electronic health record, and that they want and will trust the system's output.

We are collaborating with biomedical researchers on the design of a DST supporting the decision to implant an artificial heart. This artifact—called MATE—facilitates assist-



CHI 2019 Paper



CHI 2019, May 4–9, 2019, Glasgow, Scotland, UK

Unremarkable AI: Fitting Intelligent Decision Support into Critical, Clinical Decision-Making Processes

Qian Yang

HCI Institute

Carnegie Mellon University

yangqian@cmu.edu

Aaron Steinfeld

Robotics Institute

Carnegie Mellon University

steinfeld@cmu.edu

John Zimmerman

HCI Institute

Carnegie Mellon University

johnz@cs.cmu.edu

ABSTRACT

Clinical decision support tools (DST) promise improved health care outcomes by active in lab setting; Empirical research

This paper describes a new form of clinicians' decision support tools. This can be described as Unremarkable Continues technology, which users yet remain unaware of. Clinicians are more likely to accept DST. Drawing on the complexities of decision support systems in DST design, critical AI systems

CCS CONCEPTS
• Human-centered

KEYWORDS

1 INTRODUCTION

The idea of leveraging machine intelligence in healthcare

CHI 2019 Paper

CHI 2019, May 4–9, 2019, Glasgow, Scotland, UK

Will You Accept an Imperfect AI? Exploring Designs for Adjusting End-user Expectations of AI Systems

Rafal Kocielnik

University of Washington

Seattle, USA

rafal.kocielnik@gmail.com

Saleema Amershi

Microsoft Research

Redmond, USA

samershi@microsoft.com

Paul N. Bennett

Microsoft Research

Redmond, USA

paul.n.bennett@microsoft.com

The Scheduling Assistant can correctly detect meeting requests about 50% of the time.

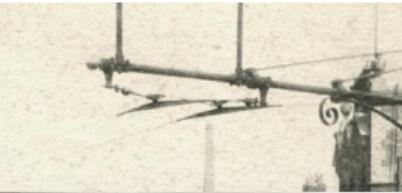
The Scheduling Assistant examines each sentence separately and looks for meeting related phrases to

Adjust how aggressive you would want the Scheduling Assistant to be in detecting meetings in your emails.





KDD 2017 Research Paper



KDD'17, August 13–17, 2017, Halifax, NS, Canada



The Selective Labels Problem: Evaluating Algorithmic Predictions in the Presence of Unobservables

Himabindu Lakkaraju
Stanford University
himalv@cs.stanford.edu

Jon Kleinberg
Cornell University
kleinber@cs.cornell.edu

Jure Leskovec
Stanford University
jure@cs.stanford.edu

Jens Ludwig
University of Chicago
jldwig@uchicago.edu

Sendhil Mullainathan
Harvard University
mullain@fas.harvard.edu

ABSTRACT

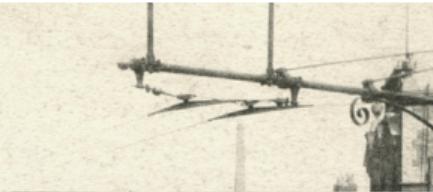
Evaluating whether machines improve on human performance is one of the central questions of machine learning. However, there are many domains where the data is *selectively labeled* in the sense that the observed outcomes are themselves a consequence of the existing choices of the human decision-makers. For instance, in the context of judicial bail decisions, we observe the outcome of whether a defendant fails to return for their court appearance only if the human judge decides to release the defendant on bail. This selective labeling makes it harder to evaluate predictive models

where the machine learning algorithm must be evaluated on data where the labels are themselves consequence of the existing choices of the human decision-makers.

We first deal with issues of this form in an analysis of judicial bail decisions [17], an application which motivated the present paper. Since this is an important setting that illustrates the basic concerns, it is useful to briefly describe the underlying background of the application here. In a bail hearing, by law requires the judge to base their decision to release defendants on a prediction—if granted bail, will the defendant return for their court appearance without



KDD 2017 Research Paper



KDD'17, August 13–17, 2017, Halifax, NS, Canada



The Selective Labels Problem: Evaluating Algorithms Predictions in the Presence

Himabindu Lakkaraju
Stanford University
himalv@cs.stanford.edu

Jon Kleinberg
Cornell University
kleinber@cs.cornell.e

Jens Ludwig
University of Chicago
jludwig@uchicago.edu

Se
mi

ABSTRACT

Evaluating whether machines improve on human performance is one of the central questions of machine learning. However, there are many domains where the data is *selectively labeled* in the sense that the observed outcomes are themselves a consequence of the existing choices of the human decision-makers. For instance, in the context of judicial bail decisions, we observe the outcome of whether a human judge fails to return for their court appearance only if the human judge decides to release the defendant on bail. This selective labeling makes it harder to evaluate predictive models

where t
data wh
choices
We fi
decisor
Since th
it is use
applical
base the
bail, wil
8
10
11
12

Human Perceptions of Fairness in Algorithmic Decision Making: A Case Study of Criminal Risk Prediction

Nina Grgić-Hlača
MPI-SWS, Saarland University
nghlaca@mpi-sws.org

Krishna P. Gummadi
MPI-SWS, Saarland University
gummadi@mpi-sws.org

ABSTRACT

Elissa M. Redmiles^{*}
University of Maryland
eredmiles@cs.umd.edu

Adrian Weller[†]
Cambridge University, Alan Turing Institute
adrian.weller@eng.cam.ac.uk

ACM Reference Format:



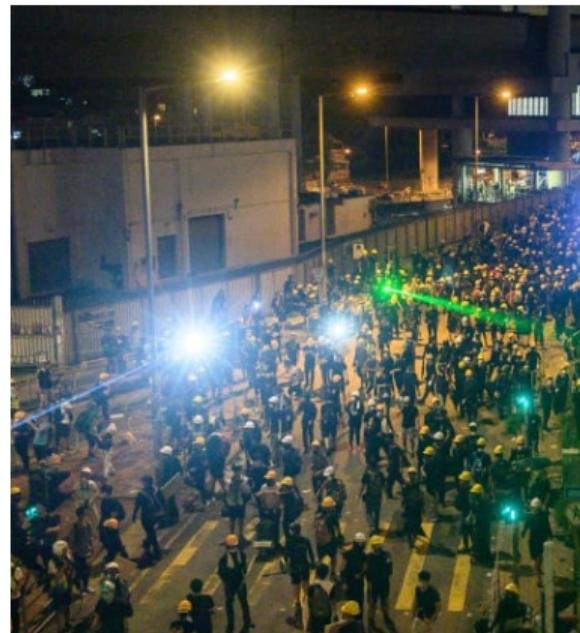
WHAT DOES A WORLD LOOK LIKE THAT IS BUILT UNCRITICALLY ON HISTORICAL FOUNDATIONS OF GREED, EXPLOITATION, AND HARM?

*how do we **perform** for algorithmic audiences right now?*

WAYS TO UNDERMINE A POLICE STATE



WAYS TO UNDERMINE A POLICE STATE



WAYS TO UNDERMINE A POLICE STATE



WHAT DOES A WORLD LOOK LIKE THAT IS BUILT UNCRITICALLY ON HISTORICAL FOUNDATIONS OF GREED, EXPLOITATION, AND HARM?

*how do we **perform** for algorithmic audiences right now?*

WHAT DOES A WORLD LOOK LIKE THAT IS BUILT UNCRITICALLY ON HISTORICAL FOUNDATIONS OF GREED, EXPLOITATION, AND HARM?

*how do we **perform** for algorithmic audiences right now?*

WHAT WOULD A WORLD LOOK LIKE THAT IS BUILT INTENTIONALLY WITH THOSE HISTORIES IN MIND?